

The results below are generated from an R script.

```
---
title: "R Notebook"
output:
  html_document:
    df_print: paged
  html_notebook: default
---

# Running on Discovery

I recommend viewing this with a web-based Rstudio server on Discovery:

https://ood.discovery.neu.edu/pun/sys/dashboard/batch_connect/sys/RStudio/session_contexts/new

Press *Ctrl+Enter* to run a chunk.

# Initialization

*You may want to change the directory below.*

```{r setup}
library(tidyverse)
knitr::opts_knit$set(root.dir = "/scratch/a.guha/minnpm-exp")
```

Load the data:

```{r}
raw_data <- read_csv("results.csv",
  col_types = cols(Status=col_factor(),
                    Project=col_factor(),
                    Rosette=col_logical(),
                    Consistency=col_factor(),
                    Minimize=col_factor(),
                    Time=col_double(),
                    NDeps=col_integer()),
  show_col_types = FALSE)
```

# Manual Verification Step

Check that these are the factors that appear below:

1. *success*: everything worked!
2. *ERESOLVE*: depends on something that isn't in the repository
3. *ETARGET*: requires some other target architecture **verify**
4. *EBADPLATFORM*: requires some other platform (e.g., macOS)
5. *EUNSUPPORTEDPROTOCOL*: a dependency is in a format that NPM does not support
6. *unexpected*: something went wrong on Discovery. See experiment.out
7. *unavailable*: something went wrong and we didn't even capture the result.
```

See experiment.out

```
```{r}
levels(raw_data$Status)
```
```

```
```{r}
levels(raw_data$Consistency)
```
```

```
```{r}
levels(raw_data$Minimize)
```
```

\*The value in the Count column should be 1000 for every row:\*

```
```{r}
raw_data %>%
  group_by(Rosette, Minimize, Consistency) %>%
  summarize(Count = n())
```
```

#### *# Failures*

How many failures occur for each configuration? We know we will see more failures than NPM, but hopefully

```
```{r}
raw_data %>%
  filter(Status != "success") %>%
  group_by(Rosette, Minimize, Consistency) %>%
  summarize(Count = n())
```
```

Let's rule out failures that are due to unsolvability on our platform:

```
```{r}
raw_data %>%
  filter(Status == "unexpected" | Status == "unavailable") %>%
  group_by(Rosette, Minimize, Consistency) %>%
  summarize(Count = n())
```
```

These are likely due to timeouts, Z3 crashing, etc.

#### *# The Need for Tree-Solving*

NPM uses a tree-solver because its easy to resolve conflicts. But, how many conflicts do you see with M

TODO: Need to process output further to distinguish these errors.

#### *# Minimizing Number of Dependencies*

For each project, the number of dependencies with vanilla NPM, and with MinNPM configured to minimize #

```

```{r}
min_dep_analysis <- bind_rows(raw_data %>%
  filter(Rosette == FALSE) %>%
  select(Project,NDeps) %>%
  mutate(Solver="NPM"),
raw_data %>%
  filter(Rosette == TRUE & Consistency == "npm" &
    Minimize == "min_num_deps,min_oldness") %>%
  select(Project, NDeps) %>%
  mutate(Solver="MinDeps")) %>%
  pivot_wider(values_from=NDeps, names_from=Solver)
```

```

<span style="color: red">

In theory, MinNPM should always produce a smaller solution. But, it seems like it doesn't always.

```

```{r}
min_dep_analysis %>% mutate(Badness = MinDeps - NPM) %>% filter(Badness > 0)
```

```

<span style="color:red">The graph below is bogus, since I've filtered out the outliers.</span>

```

```{r}
min_dep_analysis %>%
  mutate(Shrinkage = MinDeps / NPM) %>%
  filter(Shrinkage <= 1.0) %>%
  select(Shrinkage) %>%
  ggplot(aes(Shrinkage)) +
  stat_ecdf() +
  ylab("Fraction of solves that shrink")
```

```

*# Impact on Priorities*

<span style="color:red">Skip this? Not very informative.</span>

```

```{r}
bind_rows(
  raw_data %>%
    filter(Rosette == TRUE & Consistency == "npm" &
      Minimize == "min_num_deps,min_oldness") %>%
    select(Project,NDeps) %>%
    mutate(Solver="Deps,Oldness"),
raw_data %>%
  filter(Rosette == TRUE & Consistency == "npm" &
    Minimize == "min_oldness,min_duplicates") %>%
  select(Project,NDeps) %>%
  mutate(Solver="Oldness,Deps"),
raw_data %>%

```

```

  filter(Rosette == TRUE & Consistency == "npm" &
         Minimize == "min_duplicates,min_oldness") %>%
  select(Project, NDepends) %>%
  mutate(Solver="Dups,Oldness")) %>%
ggplot(aes(NDepends,color=Solver)) +
stat_ecdf()
```

```

*# Slowdown of MinNPM*

<span style="color:red">These were run on different kinds of machines, etc. and in parallel. So timing is

1. Why is it that the Rosette solver is faster? It seems almost unbelievable that the NPM solver is slow

2. Should we somehow represent the timeouts on this graph?

```

```{r}
bind_rows(
  raw_data %>%
    filter(Rosette == FALSE) %>%
    select(Project,Time) %>%
    mutate(Solver="NPM"),
  raw_data %>%
    filter(Rosette == TRUE & Consistency == "npm" &
         Minimize == "min_num_deps,min_oldness") %>%
    select(Project,Time) %>%
    mutate(Solver="MinNPM")) %>%
pivot_wider(values_from=Time, names_from=Solver) %>%
mutate(DeltaTime = MinNPM- NPM) %>%
select(-NPM, -MinNPM) %>%
ggplot(aes(y=DeltaTime)) +
stat_bin()
```

```

```

## Error: <text>:11:3: unexpected symbol
## 10:
## 11: I recommend
##      ^

```

The R session information (including the OS info, R version and all packages used):

```

sessionInfo()

## R version 4.0.3 (2020-10-10)
## Platform: x86_64-pc-linux-gnu (64-bit)
## Running under: CentOS Linux 7 (Core)
##
## Matrix products: default
## BLAS: /shared/centos7/r-project/4.0.3/lib64/R/lib/libRblas.so
## LAPACK: /shared/centos7/r-project/4.0.3/lib64/R/lib/libRlapack.so
##
## locale:
##  [1] LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C              LC_TIME=en_US.UTF-8
##  [4] LC_COLLATE=en_US.UTF-8    LC_MONETARY=en_US.UTF-8   LC_MESSAGES=en_US.UTF-8

```

```
## [7] LC_PAPER=en_US.UTF-8      LC_NAME=C          LC_ADDRESS=C
## [10] LC_TELEPHONE=C             LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods    base
##
## other attached packages:
## [1] knitr_1.36
##
## loaded via a namespace (and not attached):
## [1] compiler_4.0.3 magrittr_2.0.1 tools_4.0.3    stringi_1.7.5  highr_0.9
## [6] stringr_1.4.0  xfun_0.28     evaluate_0.14
Sys.time()
## [1] "2021-12-10 11:50:50 EST"
```